

# AZURE COGNITIVE SERVICES

Jeff Proise

# Azure Cognitive Services

- AI as a Service (AlaaS) - Deep-learning models offered thru REST APIs
- Object detection, sentiment analysis, neural machine translation (NMT), real-time speech translation, anomaly detection, and more



## Vision Services

Computer Vision  
Custom Vision  
Face\*



## Language Services

Language Service  
Translator Service



## Speech Service

Text to Speech  
Speech to Text  
Speech Translation  
Speaker Recognition



## Decision Services

Anomaly Detector  
Content Moderator  
Personalizer



## OpenAI

Managed version of  
OpenAI models and APIs  
for generative AI

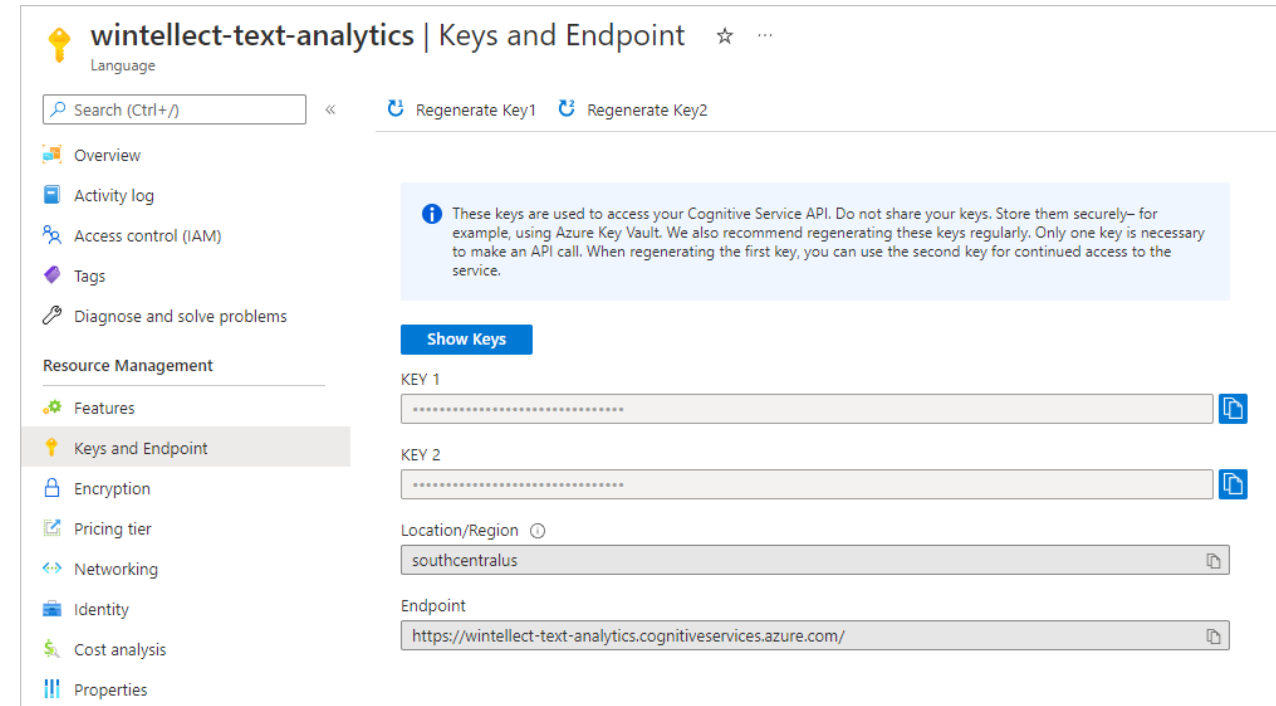
\* No longer available to the general public. Available to managed customers and partners subject to approval by Microsoft.

# Authentication

- Calls to Azure Cognitive Services must be authenticated so Microsoft can bill an Azure subscription for calls when appropriate
  - Most services have free tiers which are ideal for development
- All Cognitive Services support authentication with subscription keys obtained from the Azure Portal or the Azure CLI
  - Keys can be single-service or multi-service
- Most Cognitive Services support Azure Active Directory (AAD) authentication, too
  - Enhances security when caller is hosted in Azure

# Key Management

- Guard subscription keys as if they were passwords
- Keys should be rotated periodically for security
- Azure provides two keys for each Cognitive Services resource you deploy so you can regenerate a key without affecting apps that use the service



# Demo

Keys and Endpoints



# Calling Cognitive Services APIs

```
import requests

input = {'documents': [{'id': '1', 'text': 'Programming is fun, but the hours are long'}]}
headers = {
    'Ocp-Apim-Subscription-Key': KEY,
    'Content-type': 'application/json'
}

uri = ENDPOINT + 'text/analytics/v3.0/sentiment'
response = requests.post(uri, headers=headers, json=input)
results = response.json()

for result in results['documents']:
    print(result['confidenceScores'])
```

```
{'positive': 0.85, 'neutral': 0.03, 'negative': 0.12}
```

# Cognitive Services SDKs

- Free SDKs for a variety of programming languages simplify your code and are available for most Cognitive Services

Language (Text Analytics)	C#	Java	JavaScript	Python				
Computer Vision	C#	Java	JavaScript	Python	Go			
Speech	C#	Java	JavaScript	Python	Go	C++	Swift	Objective-C
Anomaly Detector	C#	Java	JavaScript	Python	Go			

# Using the Python Text-Analytics SDK

```
from azure.core.credentials import AzureKeyCredential
from azure.ai.textanalytics import TextAnalyticsClient

client = TextAnalyticsClient(ENDPOINT, AzureKeyCredential(KEY))
input = [{ 'id': '1', 'text': 'Programming is fun, but the hours are long' }]
response = client.analyze_sentiment(input)

for result in response:
    print(result.confidence_scores)
```

```
{'positive': 0.85, 'neutral': 0.03, 'negative': 0.12}
```



# Using the .NET Text-Analytics SDK

```
using Azure;  
using Azure.AI.TextAnalytics;  
using System;  
  
var client = new TextAnalyticsClient(new Uri(ENDPOINT), new AzureKeyCredential(KEY));  
var response = client.AnalyzeSentiment("Programming is fun, but the hours are long");  
var result = response.Value.ConfidenceScores.Positive;  
Console.WriteLine(result);
```

# Handling Errors

```
from azure.core.credentials import AzureKeyCredential
from azure.ai.textanalytics import TextAnalyticsClient
from azure.core.exceptions import AzureError

try:
    client = TextAnalyticsClient(ENDPOINT, AzureKeyCredential(KEY))
    input = [{ 'id': '1', 'text': 'Programming is fun, but the hours are long' }]
    response = client.analyze_sentiment(input)

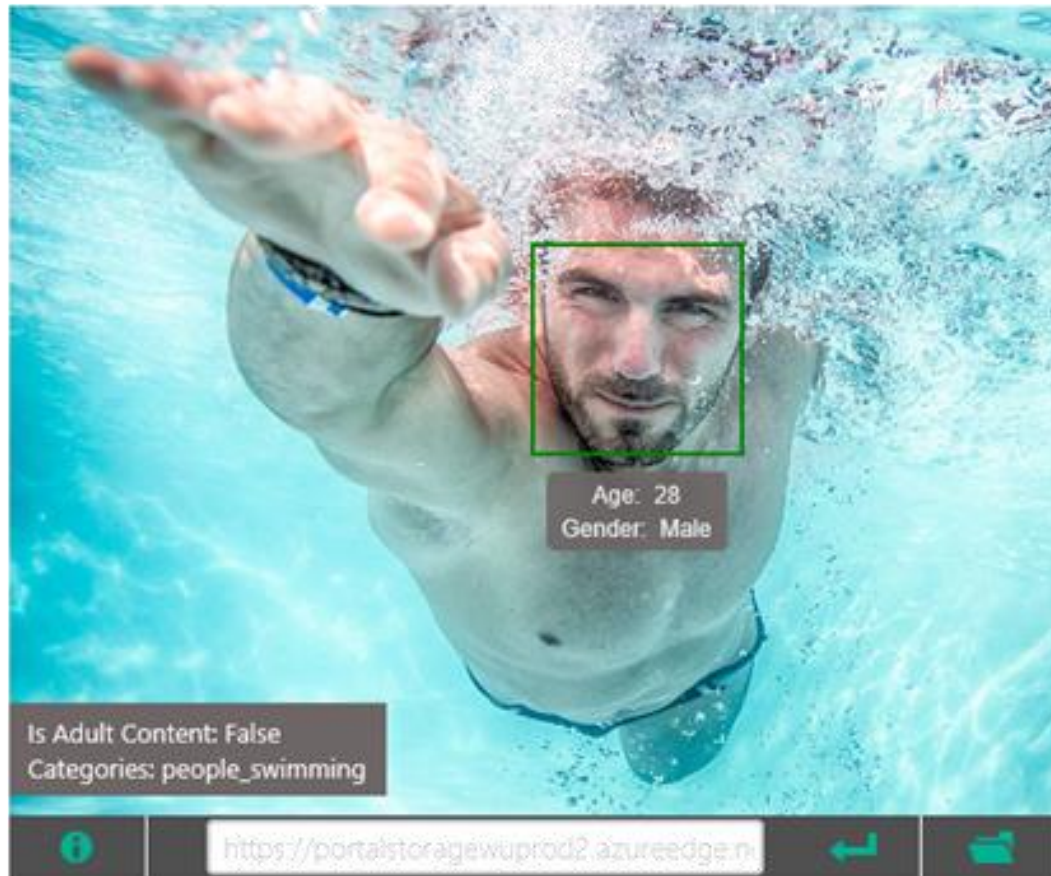
    for result in response:
        print(result.confidence_scores)

except AzureError as e:
    print(e.message)
```

# Computer Vision Service

- Analyze images and extract information from them
  - Caption photos
  - Generate tags describing the contents of a photo
  - Detect and identify objects in photos (including bounding boxes)
  - Detect faces in photos and identify ages and genders
  - Identify images containing adult or inappropriate content
  - Generate "smart thumbnails," extract text from images, and more
- Python SDK in **azure-cognitiveservices-vision-computervision**
- **ComputerVisionClient** class provides wrapper around REST API

# Computer Vision Service in Action



## Features:

### Feature Name

### Value

### Description

{ "type": 0, "captions": [ { "text": "a man swimming in a pool of water", "confidence": 0.7850108693093019 } ] }

### Tags

[ { "name": "water", "confidence": 0.9996442794799805 }, { "name": "sport", "confidence": 0.9504992365837097 }, { "name": "swimming", "confidence": 0.9062818288803101, "hint": "sport" }, { "name": "pool", "confidence": 0.8787588477134705 }, { "name": "water sport", "confidence": 0.631849467754364, "hint": "sport" } ]

### Image Format

jpeg

### Image Dimensions

1500 x 1155

### Clip Art Type

0 Non-clipart

### Line Drawing Type

0 Non-LineDrawing

### Black & White Image

False

# Captioning an Image

```
client = ComputerVisionClient(ENDPOINT, CognitiveServicesCredentials(KEY))

# Submit image from local file system
with open('IMAGE_PATH', mode='rb') as image:
    result = client.describe_image_in_stream(image)
    for caption in result.captions:
        print(f'{caption.text} ({caption.confidence:.1%})')

# Submit image URL
result = client.describe_image('IMAGE_URL')
for caption in result.captions:
    print(f'{caption.text} ({caption.confidence:.1%})')
```

# Generating Keywords from Images

```
client = ComputerVisionClient(ENDPOINT, CognitiveServicesCredentials(KEY))

# Submit image from local file system
with open('IMAGE_PATH', mode='rb') as image:
    result = client.tag_image_in_stream(image)
    for tag in result.tags:
        print(f'{tag.name} ({tag.confidence:.1%})')
```

# Detecting Objects in Images

```
client = ComputerVisionClient(ENDPOINT, CognitiveServicesCredentials(KEY))

# Submit image from local file system
with open('IMAGE_PATH', mode='rb') as image:
    result = client.detect_objects_in_stream(image)
    for obj in result.objects:
        print(f'{obj.object_property}, {obj.confidence:.1%}, {obj.rectangle}')
```



*Object name (person, dog, bicycle, etc.)*



*Bounding box*

# Demo

The Computer Vision Service





# Language Service

- Supports text analytics (e.g., sentiment analysis, named-entity recognition, and key-phrase extraction), question answering (semantic search), conversational language understanding, and more
- Helper classes available in Python SDKs
  - **TextAnalyticsClient** class in **azure-ai-textanalytics**
  - **QuestionAnsweringClient** class in **azure-ai-language-questionanswering**
  - **ConversationAnalysisClient** class in **azure-ai-language-conversations**
- Used by LaLiga to boost fan engagement with conversational AI

# Analyzing Sentiment

```
client = TextAnalyticsClient(ENDPOINT, AzureKeyCredential(KEY))
input = [{ 'id': '1', 'text': 'Programming is fun, but the hours are long' }]
response = client.analyze_sentiment(input)

for result in response:
    print(result.confidence_scores)
```

```
{'positive': 0.85, 'neutral': 0.03, 'negative': 0.12}
```

# Analyzing Sentiment in Batches

```
input = [  
    { 'id': '1', 'text': 'Programming is fun, but the hours are long' },  
    { 'id': '2', 'text': 'Great food and excellent service' },  
    { 'id': '3', 'text': 'The product worked as advertised but is overpriced' },  
    { 'id': '4', 'text': 'Moving to the cloud was the best decision we ever made' }  
]
```

```
client = TextAnalyticsClient(ENDPOINT, AzureKeyCredential(KEY))  
response = client.analyze_sentiment(input)
```

```
for result in response:  
    text = ''.join([x.text for x in result.sentences])  
    print(f'{text} => {result.confidence_scores.positive}')
```

# Question Answering

```
from azure.ai.language.questionanswering import models as qna
from azure.ai.language.questionanswering import QuestionAnsweringClient

client = QuestionAnsweringClient(ENDPOINT, AzureKeyCredential(KEY))
question = 'What is the minimum age required to serve as a United State Senator?'

context = 'No Person shall be a Senator who shall not have attained to the Age of ' \
          'thirty Years, and been nine Years a Citizen of the United States, and ' \
          'who shall not, when elected, be an Inhabitant of that State for which ' \
          'he shall be chosen.'

input = qna.AnswersFromTextOptions(question=question, text_documents=[context])
results = client.get_answers_from_text(input)
answer = results.answers[0].short_answer
print(f'{answer.text} ({answer.confidence:.1%})')
```

thirty Years (88.0%)

# Translator Service

- Uses neural machine translation (NMT) to translate text to other languages and understands >100 written languages and dialects
- Supports two APIs (one service, two endpoints)
  - Text translation API - No Python SDK available at this time
  - Document translation API - SDK available in [azure-ai-translation-document](#) to batch-translate documents (including PDFs) at scale in Azure Blob Storage
- Used by Volkswagen to translate onscreen instruction in cars and translate documentation into more than 40 languages

# Translating French to English

```
input = [{ 'text': 'Quand votre nouveau livre sera-t-il disponible?' }]
```

```
headers = {  
    'Ocp-Apim-Subscription-Key': KEY,  
    'Ocp-Apim-Subscription-Region': REGION,  
    'Content-type': 'application/json'  
}
```

```
uri = ENDPOINT + 'translate?api-version=3.0&from=fr&to=en'  
response = requests.post(uri, headers=headers, json=input)  
results = response.json()
```

```
print(results[0]['translations'][0]['text'])
```

When will your new book be available?

# Demo

Language Services



# Speech Service

- Converts text to speech and speech to text, identifies spoken languages, and translates speech to other languages in real time
- Python SDK available in **azure-cognitiveservices-speech**
  - **SpeechRecognizer** class converts speech to text
  - **SpeechSynthesizer** class converts text to speech in more than 300 voices
  - **TranslationRecognizer** class converts speech to text and translates it, too
- Used by Airbus to build voice-enabled apps for pilots and by KPMG to transcribe calls and reduce compliance costs for clients



# Converting Text to Speech

```
from azure.cognitiveservices import speech
```

```
config = speech.SpeechConfig(KEY, REGION)
```

```
config.speech_synthesis_voice_name = 'en-US-JennyNeural'
```

```
synthesizer = speech.SpeechSynthesizer(config)
```

```
synthesizer.speak_text_async('When will your new book be published?').get()
```

```
# For a complete list of more than 300 neural voices, see https://docs.microsoft.com/en-us/azure/cognitive-services/speech-service/language-support?tabs=speechtotext#prebuilt-neural-voices
```

# Demo

The Speech Service



# OpenAI

- Creator of ChatGPT, DALL·E 2, and other cutting-edge models
- Deploy models to Azure and access REST APIs for generating text, images, code, text embeddings, and more

**GPT-3**  
**GPT-3.5**  
**GPT-4**

Models that support **text and code generation**. Includes **ChatGPT**.

**Codex**

Models that support **code generation**, converting code to other languages, and more. **Discontinued on March 23, 2023**.

**DALL·E 2**

Produces **images from natural-language prompts**. Supports inpainting, outpainting, image variations, and more.

**Embeddings**

Models that generate **embedding vectors** from text. Used for semantic search, recommender systems, and more.

**Whisper**

General-purpose **speech-to-text engine** trained on 680,000 hours of audio. Also available via open source.

# Generating Text with ChatGPT

```
import openai
openai.api_type = 'azure'
openai.api_version = '2023-03-15-preview'
openai.api_base = 'API_ENDPOINT' # Obtained from Azure portal
openai.api_key = 'API_KEY' # Obtained from Azure portal

messages = [{ 'role': 'user', 'content': 'Write a poem about deep learning' }]

response = openai.ChatCompletion.create(
    engine='my-chatgpt-instance', # Azure deployment name
    messages=messages
)
```

# Streaming the Response

```
messages = [{ 'role': 'user', 'content': 'Write a poem about deep learning' }]

chunks = openai.ChatCompletion.create(
    engine='my-chatgpt-instance',
    messages=messages,
    stream=True
)

for chunk in chunks:
    content = chunk['choices'][0].get('delta', {}).get('content')
    if content is not None:
        print(content, end='')

```

# Translating Text

```
content = 'Translate the following text from English to French: Best food ever!'
```

```
messages = [{ 'role': 'user', 'content': content }]
```

```
response = openai.ChatCompletion.create(  
    engine='my-chatgpt-instance',  
    messages=messages,  
    temperature=0  
)
```

# Generating Code

```
content = 'Create a Python function that accepts an array of numbers as ' \
          'input, bubble sorts the numbers, and returns a sorted array'

messages = [{ 'role': 'user', 'content': content }]

response = openai.ChatCompletion.create(
    engine='my-chatgpt-instance',
    messages=messages,
    temperature=0
)
```

# Handling Errors

```
try:
    response = openai.ChatCompletion.create(
        engine='my-chatgpt-instance',
        messages=messages
    )

    print(response.choices[0].message.content)

except AuthenticationError as e:
    print('Invalid API key')
except ServiceUnavailableError as e:
    print('ChatGPT temporarily unavailable')
except Exception as e:
    print('Call failed')
```



atmosera®



# Demo

OpenAI



# Custom Vision Service

- Build custom image-classification and object-detection models
  - Get acceptable accuracy with as few as 50 to 100 training images
- Build intelligent apps that invoke models using REST API
- Or export to CoreML or TensorFlow and run models locally

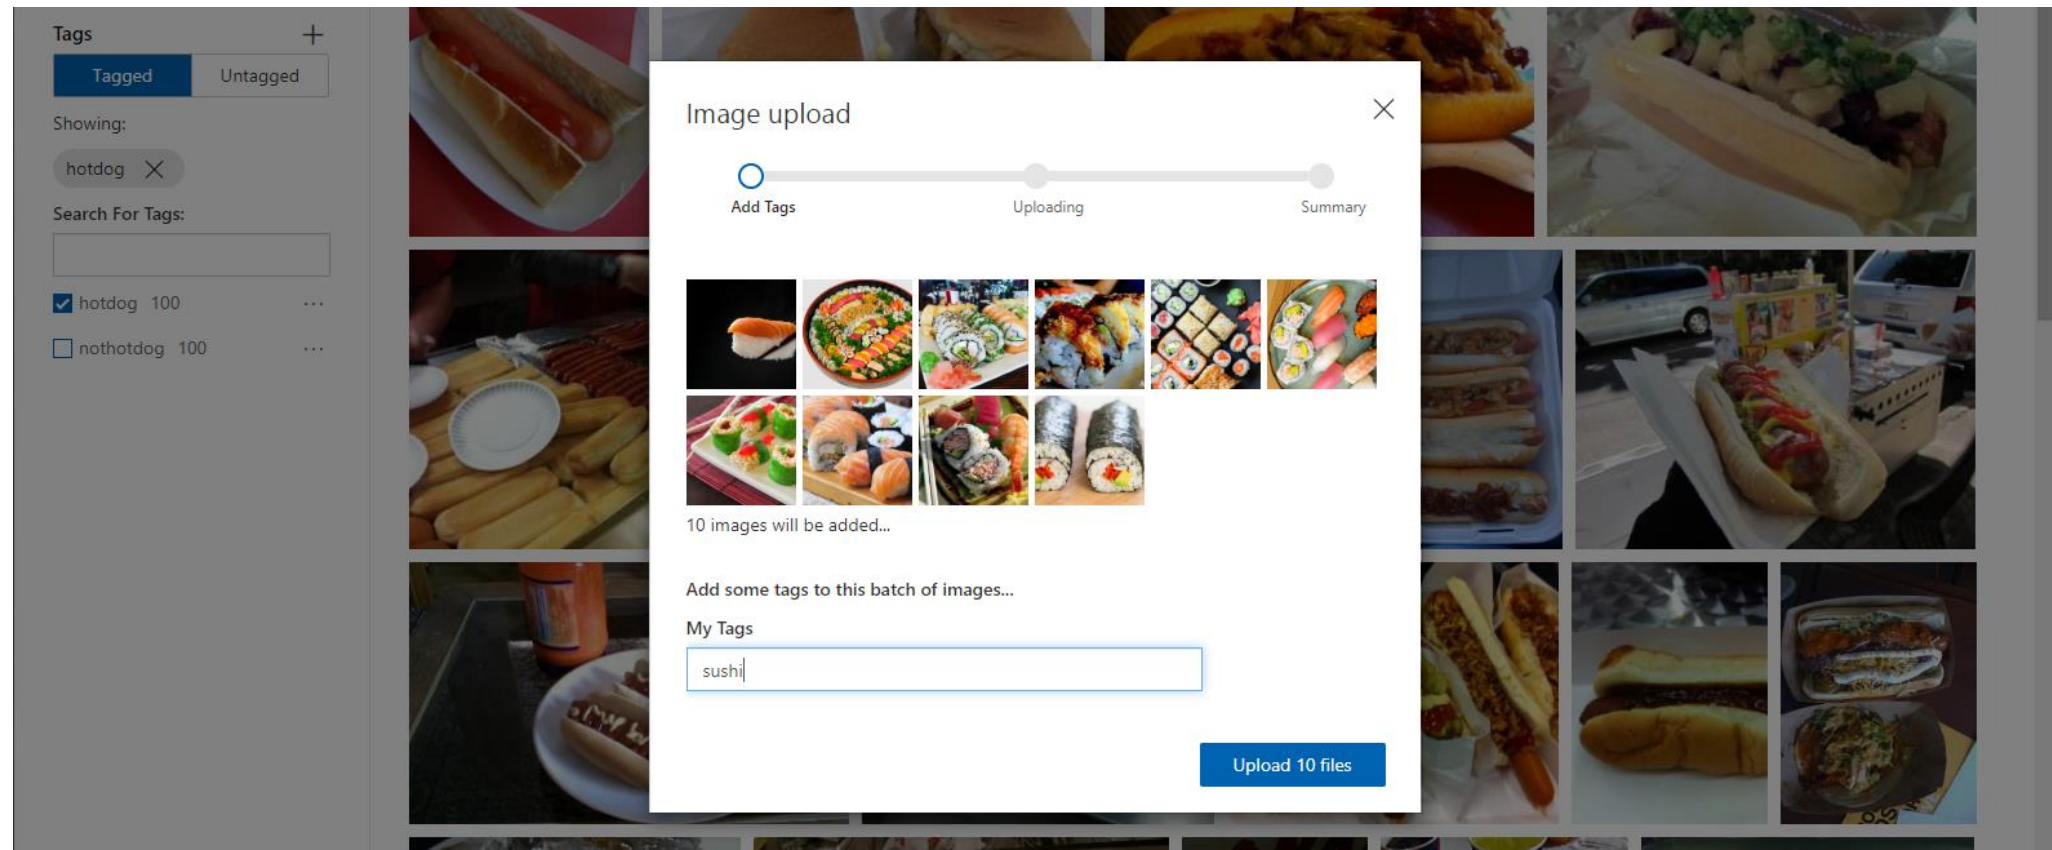


## Results

Tag	Probability
daisy	99.9%
trillium	3.1%
lily of the valley	0.1%
dogwood	0.0%


# Training a Custom Vision Service Model

- Use Custom Vision portal ([www.customvision.ai](https://www.customvision.ai)) to create projects, upload and tag images, and train and publish models



# Demo

The Custom Vision Service



# Cognitive Services Containers

- Most Cognitive Services can be hosted locally in Docker containers
  - Keep data on-premises (avoid sending it to Azure)
  - Insulate your code from API changes and updates
  - Run Cognitive Services physically close to apps that use them
  - Use Kubernetes or other container orchestration platforms to control scaling
- Containerized services do NOT prevent you from being billed
  - Containers send metering information to Microsoft using port 443

# Disconnected Containers

- Containers that run without an Internet connection
  - Do not send metering information to Microsoft
  - For environments where connectivity is non-existent or spotty or policies and regulations forbid apps that use Cognitive Services from accessing the Internet
- Only available for subsets of certain services (e.g., text extraction via Computer Vision service)
- Only available to "strategic customers and partners" with Enterprise Agreements (EAs) in place and must be approved by Microsoft